

You are not subscribed. [Learn more](#)

Available: 58.41 compute units

Usage rate: approximately 4.82 per hour

You have 1 active session.

[Manage sessions](#)

Want more memory and disk space?



[Upgrade to Colab Pro](#)

Not connected to runtime.

[Change runtime type](#)

```
%cd /content/drive/MyDrive/ImageBind finetune/ImageBind-LoRA
```

```
↗ /content/drive/MyDrive/ImageBind finetune/ImageBind-LoRA
```

```
!pip install -r requirements.txt --quiet
!pip install pytorch_lightning --quiet
```

```
↗ Preparing metadata (setup.py) ... done
_____ 890.1/890.1 MB 1.3 MB
_____ 24.3/24.3 MB 66.0 MB/
_____ 4.2/4.2 MB 98.3 MB/s
_____ 510.0/510.0 kB 51.8 M
_____ 54.4/54.4 kB 8.5 MB/s
_____ 43.2/43.2 kB 5.9 MB/s
_____ 50.2/50.2 kB 7.4 MB/s

Preparing metadata (setup.py) ... done
_____ 13.6/13.6 MB 98.6 MB/
_____ 42.2/42.2 kB 6.0 MB/s

Preparing metadata (setup.py) ... done
_____ 7.1/7.1 MB 119.0 MB/s

Installing build dependencies ... done
Getting requirements to build wheel ... done
Preparing metadata (pyproject.toml) ... done
_____ 11.6/11.6 MB 119.0 MB
_____ 1.8/1.8 MB 91.0 MB/s
_____ 849.3/849.3 kB 65.8 M
_____ 557.1/557.1 MB 2.1 MB
_____ 317.1/317.1 MB 2.7 MB
_____ 21.0/21.0 MB 80.8 MB/
_____ 66.4/66.4 kB 9.3 MB/s
_____ 80.8/80.8 kB 12.0 MB/
_____ 55.5/55.5 kB 8.2 MB/s
_____ 1.2/1.2 MB 58.1 MB/s
_____ 71.9/71.9 kB 10.3 MB/
_____ 868.8/868.8 kB 59.2 M
_____ 62.4/62.4 kB 10.5 MB/
_____ 129.9/129.9 kB 18.5 M
_____ 812.3/812.3 kB 61.1 M
_____ 34.3/34.3 MB 50.9 MB/
_____ 230.0/230.0 kB 29.1 M
_____ 268.9/268.9 kB 33.0 M
_____ 1.3/1.3 MB 74.5 MB/s
_____ 5.1/5.1 MB 110.4 MB/s
_____ 1.5/1.5 MB 84.9 MB/s
_____ 3.1/3.1 MB 99.4 MB/s
_____ 64.3/64.3 kB 9.6 MB/s
_____ 58.4/58.4 kB 8.6 MB/s
_____ 139.2/139.2 kB 19.2 M
_____ 58.3/58.3 kB 8.2 MB/s
_____ 812.3/812.3 kB 62.2 M
_____ 812.3/812.3 kB 63.6 M
_____ 812.2/812.2 kB 63.4 M
_____ 802.3/802.3 kB 63.7 M
_____ 12.4/12.4 MB 106.2 MB
_____ 82.7/82.7 kB 12.9 MB/

Building wheel for pytorchvideo (setup.py) ... done
Building wheel for fvcore (setup.py) ... done
Building wheel for iopath (setup.py) ... done
Building wheel for mayavi (pyproject.toml) ... done
ERROR: pip's dependency resolver does not currently take into account
torchtext 0.18.0 requires torch>=2.3.0, but you have torch 1.13.0 w
```

◀ ▶

```
import logging
logging.basicConfig(level=logging.INFO, force=True)
```

```
import os
num_workers = os.cpu_count()
print(f"Number of CPU cores: {num_workers}")
```

```
↗ Number of CPU cores: 12
```

```
from pytorch_lightning import seed_everything
seed_everything(43, workers=True)
```



```
INFO:lightning_fabric.utilities.seed:Seed set to 43
43
```

```
import os
from torch.utils.data import Dataset
from sklearn.model_selection import train_test_split
```

```
import data
```

```
/usr/local/lib/python3.10/dist-packages/torchvision/transforms/_fun
warnings.warn(
/usr/local/lib/python3.10/dist-packages/torchvision/transforms/_tra
warnings.warn(
```

```
import torch
import torch.nn.functional as F
import torch.optim as optim
from torch.utils.data import DataLoader, ConcatDataset
import torchvision
from torchvision import transforms

from models import imagebind_model
from models import lora as LoRA
from models.imagebind_model import ModalityType, load_module, save_modu
```

```
import pytorch_lightning as L
from pytorch_lightning import Trainer, seed_everything
from pytorch_lightning.callbacks import ModelCheckpoint
from pytorch_lightning import loggers as pl_loggers
```

```
self_contrast = False
batch_size = 4
num_workers= os.cpu_count()
lora_modality_names_123 = ["vision", "audio"]
LOG_ON_STEP = False
LOG_ON_EPOCH = True
lora= True
full_model_checkpointing = False
full_model_checkpoint_dir="./checkpoints/full"
lora_checkpoint_dir="./checkpoints/lora"
device_name="cuda:0" if torch.cuda.is_available() else "cpu"
max_epochs = 5
gradient_clip_val=1.0
loggers = None
linear_probing = False
```

```
class ImageBindTrain(L.LightningModule):
    def __init__(self, lr=5e-4, weight_decay=1e-4, max_epochs=500, batch_
        self_contrast=False, temperature=0.07, momentum_betas=(
        lora=False, lora_rank=4, lora_checkpoint_dir="./checkpo
        lora_layer_idxs=None, lora_modality_names=None,
        linear_probing=False
    ):
        super().__init__()
        assert not (linear_probing and lora), \
            "Linear probing is a subset of LoRA training procedure for Im
            "Cannot set both linear_probing=True and lora=True. " \
            "Linear probing stores params in lora_checkpoint_dir"
        self.save_hyperparameters()

    # Load full pretrained ImageBind model
    self.model = imagebind_model.imagebind_huge(pretrained=True)
    if lora:
        for modality_preprocessor in self.model.modality_preprocessor
            modality_preprocessor.requires_grad_(False)
        for modality_trunk in self.model.modality_trunks.children():
            modality_trunk.requires_grad_(False)

        self.model.modality_trunks.update(LoRA.apply_lora_modality_tr
```



```

        LoRA.load_lora_modality_trunks(self.model.modality_trunks, ch

    # Load postprocessors & heads
    load_module(self.model.modality_postprocessors, module_name="
                checkpoint_dir=lora_checkpoint_dir)
    load_module(self.model.modality_heads, module_name="heads",
                checkpoint_dir=lora_checkpoint_dir)
elif linear_probing:
    for modality_preprocessor in self.model.modality_preprocessor
        modality_preprocessor.requires_grad_(False)
    for modality_trunk in self.model.modality_trunks.children():
        modality_trunk.requires_grad_(False)
    for modality_postprocessor in self.model.modality_postprocess
        modality_postprocessor.requires_grad_(False)

    load_module(self.model.modality_heads, module_name="heads",
                checkpoint_dir=lora_checkpoint_dir)
    for modality_head in self.model.modality_heads.children():
        modality_head.requires_grad_(False)
        final_layer = list(modality_head.children())[-1]
        final_layer.requires_grad_(True)

def configure_optimizers(self):
    optimizer = optim.AdamW(self.parameters(), lr=self.hparams.lr, we
        betas=self.hparams.momentum_betas)
    lr_scheduler = optim.lr_scheduler.CosineAnnealingLR(
        optimizer, T_max=self.hparams.max_epochs, eta_min=self.hparam
    )
    return [optimizer], [lr_scheduler]

def info_nce_loss(self, batch, mode="train"):
    data_a, class_a, data_b, class_b = batch

    # class_a is always "vision" according to ImageBind
    feats_a = [self.model({class_a[0]: data_a_i}) for data_a_i in dat
    feats_a_tensor = torch.cat([list(dict_.values())[0] for dict_ in
    # class_b could be any modality
    feats_b = [self.model({class_b[idx]: data_b_i}) for idx, data_b_i
    feats_b_tensor = torch.cat([list(dict_.values())[0] for dict_ in
    temperatures = [self.hparams.temperature]
    # print("feats_a_tensor.shape", feats_a_tensor.shape)
    # print("feats_b_tensor.shape", feats_b_tensor.shape)

    feats_a_normalized = F.normalize(feats_a_tensor, p=2, dim=1)
    feats_b_normalized = F.normalize(feats_b_tensor, p=2, dim=1)

    # Compute cosine similarity matrix
    cos_sim = torch.mm(feats_a_normalized, feats_b_normalized.t())

    #cos_sim = F.cosine_similarity(feats_a_tensor, feats_b_tensor.T,

    nll= False
    for i in range(cos_sim.shape[0]):
        pos_mask = [[False for _ in range(cos_sim.shape[0])] for _ in r
        for j in range(cos_sim.shape[0]):
            if i == j:
                pos_mask[i][j] = True
            else:
                pos_mask[i][j] = False
        pos_mask = torch.tensor(pos_mask, dtype=torch.bool)
        neg_mask = ~ pos_mask
        # print("pos_mask.shape", pos_mask.shape)
        # print("neg_mask.shape", neg_mask.shape)
        # print("cos_sim.shape", cos_sim.shape)
        # print("cos_sim[pos_mask].shape", cos_sim[pos_mask].shape)
        # print("cos_sim[neg_mask].shape", cos_sim[neg_mask].shape)

        if not nll:
            nll = -cos_sim[pos_mask] + torch.logsumexp(cos_sim[neg_mask],
        else:
            nll += -cos_sim[pos_mask] + torch.logsumexp(cos_sim[neg_mask]
    nll=nll.mean()
    self.log(mode + "_loss", nll, prog_bar=True,
        on step=LOG ON STEP. on epoch=LOG ON EPOCH. batch size=

```



```
        return nll

    def training_step(self, batch, batch_idx):
        return self.info_nce_loss(batch, mode="train")

    def validation_step(self, batch, batch_idx):
        self.info_nce_loss(batch, mode="val")

    def on_validation_epoch_end(self):
        if self.hparams.lora:
            # Save LoRA checkpoint
            LoRA.save_lora_modality_trunks(self.model.modality_trunks, ch
            # Save postprocessors & heads
            save_module(self.model.modality_postprocessors, module_name="
                        checkpoint_dir=self.hparams.lora_checkpoint_dir)
            save_module(self.model.modality_heads, module_name="heads",
                        checkpoint_dir=self.hparams.lora_checkpoint_dir)
        elif self.hparams.linear_probing:
            # Save postprocessors & heads
            save_module(self.model.modality_heads, module_name="heads",
                        checkpoint_dir=self.hparams.lora_checkpoint_dir)
```



```

class ImageAudioDataset(Dataset):
    def __init__(self, root_dir, transform=None, split='train', train_s:
        self.root_dir = root_dir
        self.transform = transform
        self.device = device

    self.classes = [d for d in os.listdir(os.path.join(root_dir, 'i
    self.class_to_idx = {cls: idx for idx, cls in enumerate(self.cls

    self.image_paths = []
    self.audio_paths = []
    for cls in self.classes:
        cls_image_dir = os.path.join(root_dir, 'images', cls)
        cls_audio_dir = os.path.join(root_dir, 'audio', cls)
        for filename in os.listdir(cls_image_dir):
            filename_temp=filename[:-4]
            if filename_temp[:-4] == ".DS_S":
                continue
            self.image_paths.append((os.path.join(cls_image_dir, fi
            self.audio_paths.append((os.path.join(cls_audio_dir, fi

    # Split dataset
    self.train_image_paths, self.test_image_paths = train_test_split
    self.train_audio_paths, self.test_audio_paths = train_test_split

    if split == 'train':
        self.image_paths = self.train_image_paths
        self.audio_paths = self.train_audio_paths
    elif split == 'test':
        self.image_paths = self.test_image_paths
        self.audio_paths = self.test_audio_paths
    else:
        raise ValueError(f"Invalid split argument. Expected 'train'

    def __len__(self):
        return min(len(self.image_paths), len(self.audio_paths))

    def __getitem__(self, index):
        img_path, class_text = self.image_paths[index]
        audio_path, _ = self.audio_paths[index]
        # Load and transform image
        images = data.load_and_transform_vision_data([img_path], self.d
        if self.transform is not None:
            image = images[0]
            images = self.transform(image)

        # Load and transform audio
        audios = data.load_and_transform_audio_data([audio_path], self.

        # Load and transform text
        texts = data.load_and_transform_text([class_text], self.device)

        return images, ModalityType.VISION, audios, ModalityType.AUDIO#

```



```

contrast_transforms = transforms.Compose(
    [
        transforms.RandomHorizontalFlip(),
        transforms.RandomResizedCrop(size=224),
        transforms.RandomApply([transforms.ColorJitter(brightness=0
                                                    p=0.8),
                                transforms.RandomGrayscale(p=0.2),
                                transforms.GaussianBlur(kernel_size=9),
                                transforms.ToTensor(),
                                transforms.Normalize(
                                    mean=(0.48145466, 0.4578275, 0.40821073),
                                    std=(0.26862954, 0.26130258, 0.27577711),
                                )],
            p=0.8),
    ])

class ContrastiveTransformations:
    def __init__(self, base_transforms, n_views=2):
        self.base_transforms = base_transforms
        self.n_views = n_views

    def __call__(self, x):
        return [self.base_transforms(x) for _ in range(self.n_views)]

```

```

train_datasets = []
test_datasets = []

```

```

train_datasets.append(ImageAudioDataset(
    root_dir=os.getcwd()+"/new_data/", split="train",
    transform=ContrastiveTransformations(contrast_transforms,
                                          n_views=2 if self_conti

```

```

test_datasets.append(ImageAudioDataset(
    root_dir=os.getcwd()+"/new_data/", split="test",
    transform=ContrastiveTransformations(contrast_transforms,
                                          n_views=2 if self_conti

```

```

train_dataset = train_datasets[0]
test_dataset = test_datasets[0]

```

```

train_loader = DataLoader(
    train_dataset,
    batch_size=batch_size,
    shuffle=True,
    drop_last=True,
    pin_memory=False,
    num_workers=num_workers,
)
val_loader = DataLoader(
    test_dataset,
    batch_size=batch_size,
    shuffle=False,
    drop_last=False,
    pin_memory=False,
    num_workers=num_workers,
)

```

```

lora_layer_idxs = {}
lora_modality_names = []
modalities = ["vision", "text", "audio", "thermal", "depth", "imu"]
for modality_name in lora_modality_names_123:
    if modality_name in modalities:
        modality_type = getattr(ModalityType, modality_name.upper())
        #lora_layer_idxs[modality_type] = getattr(args, f'lora_layer_id:
        # if not lora_layer_idxs[modality_type]:
        #     lora_layer_idxs[modality_type] = None
        lora_layer_idxs[modality_type] = None
        lora_modality_names.append(modality_type)
    else:
        raise ValueError(f"Unknown modality name: {modality_name}")

```




```

model = ImageBindTrain(
    max_epochs=max_epochs, batch_size=batch_size,
    num_workers=num_workers, self_contrast=self_contrast,
    lora=lora, lora_checkpoint_dir=lora_checkpoint_dir,
    lora_layer_idx=loralayer_idx if loralayer_idx is not None else None,
    lora_modality_names=lora_modality_names if lora_modality_names is not None else None,
    linear_probing=linear_probing
)

```

⚠ WARNING:root:Could not find LoRA parameters for modality vision in
 WARNING:root:If you are training the sub-model from scratch, this is
 WARNING:root:If you are loading parts of a pre-trained model, this is
 WARNING:root:Could not find LoRA parameters for modality audio in .
 WARNING:root:If you are training the sub-model from scratch, this is
 WARNING:root:If you are loading parts of a pre-trained model, this is
 WARNING:root:Could not load module parameters for postprocessors fr
 WARNING:root:Could not load module parameters for heads from ./che



```

if full_model_checkpointing:
    checkpointing = {"enable_checkpointing": full_model_checkpointing,
                    "callbacks": [ModelCheckpoint(monitor="val_loss",
                                                  filename="imagebind_checkpoint",
                                                  save_last=True)
                                ]}

```

